

## APPENDIX

### A. Proof of Proposition 1

As shown in Algorithm 2, the first step (Line 1) constructs distance matrices for all attributes by calling the algorithm PDM. For each attribute  $A_j$ , PDM divides the rows of the distance matrix among the  $|\mathcal{P}|$  processors so that each processor computes the distances for all tuple pairs in its assigned rows. After combining partial results, this yields the entire distance matrix. Since there are  $m$  attributes and  $n \times n$  tuple pairs, the overall cost is  $O\left(\frac{m \cdot n^2}{|\mathcal{P}|}\right)$ .

Next, Algorithm 2 invokes TAGC (Line 3) to generate clusters for each attribute, as presented in Algorithm 1. In particular, Line 1 sorts the thresholds  $\Theta_j$  in constant time, because the size of  $\Theta_j$  is small, while Line 2 must traverse the entire distance matrix in  $O(n^2)$  time to obtain  $C_j[\theta_j^{\max}]$ . Then, in Lines 3–5, TAGC refines that maximum-threshold clustering without a second full traversal, instead iterating over the previously computed clusters and thus costing less than  $O(n^2)$ . Throughout this process, the distance matrix rows are partitioned among the  $|\mathcal{P}|$  processors, so that each processor handles clustering for all relevant thresholds  $\Theta_j$  within its assigned rows. Consequently, generating clusters for every attribute in parallel over  $m$  attributes requires  $O\left(\frac{m \cdot n^2}{|\mathcal{P}|}\right)$  time overall.

Afterward, in Lines 4–16, Algorithm 2 conducts the complete error detection process for every constraint  $\varphi \in \Sigma$ . Specifically, in Line 6, the algorithm ASPT partitions the tuple pair sets from all relevant attribute clusters of  $\varphi$  across  $|\mathcal{P}|$  processors in parallel. Since these tuple pair sets are assigned by rows and the distance matrix is of size  $n \times n$  with  $m$  attributes, ASPT takes  $O\left(\frac{m \cdot n}{|\mathcal{P}|}\right)$  time. Next, in Line 7, Algorithm 2 calls PFVP to find violating tuple pairs by performing set operations (intersection and difference) among the tuple pair sets derived from each attribute cluster of  $\varphi$ . Because each cluster can contain up to  $O(n^2)$  tuple pairs, the row-based partitioning strategy also applies here: each processor works on the subset of rows assigned to it, carrying out intersection and difference in parallel. As a result, executing these set operations for  $m$  attributes among  $|\mathcal{P}|$  processors costs  $O\left(\frac{m \cdot n^2}{|\mathcal{P}|}\right)$ . Having obtained all violating tuple pairs  $C_\varphi^p$ , Lines 8–16 traverse each pair  $(t_i, t_j)$  to identify erroneous cells in constant time, thus taking  $O(n^2)$  time overall in that final step.

Summing up, for each constraint  $\varphi \in \Sigma$ , the cost is  $O\left(\frac{m \cdot n^2}{|\mathcal{P}|}\right)$ . Because there are  $|\Sigma|$  such constraints, the total time complexity of PED is  $O\left(\frac{|\Sigma| \cdot m \cdot n^2}{|\mathcal{P}|}\right)$  as stated in proposition 1.

### B. Proof of Proposition 2

Algorithm IPED begins by calling IPDM, which incrementally updates the distance matrix from size  $n \times n$  to  $(n + \Delta n) \times (n + \Delta n)$ . Since only rows  $1 \dots (n + \Delta n)$  and columns  $n \dots (n + \Delta n)$  must be filled, each of the  $|\mathcal{P}|$  processors is assigned a subset of these rows and computes the tuple pair distances within them, resulting in a total cost of  $O\left(\frac{m \cdot (n + \Delta n) \cdot \Delta n}{|\mathcal{P}|}\right)$  for  $m$  attributes.

Next, the algorithm IPED invokes TAGC to build incremental clusters  $\Delta C_j$  for each attribute  $A_j$  based on the new incremental distance matrix  $\Delta D_j$ . TAGC partitions the rows  $1 \dots (n + \Delta n)$  among

the  $|\mathcal{P}|$  processors, generating the clusters for its assigned rows. Because the relevant tuple pairs now total  $(n + \Delta n) \cdot \Delta n$  instead of  $n^2$ , the cost to generate clusters for  $m$  attributes is  $O\left(\frac{m \cdot (n + \Delta n) \cdot \Delta n}{|\mathcal{P}|}\right)$ .

Afterward, for each constraint  $\varphi \in \Sigma$ , the algorithm IPED detects errors incrementally. First, ASPT distributes the tuple pair sets derived from the clusters of  $\varphi$  among the  $|\mathcal{P}|$  processors by rows, so traversing  $(n + \Delta n)$  rows with  $m$  attributes costs  $O\left(\frac{m \cdot (n + \Delta n)}{|\mathcal{P}|}\right)$ . Then, PFVP finds violating tuple pairs by set operations (intersection and difference) over these tuple pair sets. Each cluster contains up to  $(n + \Delta n) \cdot \Delta n$  tuples pairs, so carrying out these operations for  $m$  attributes in parallel takes  $O\left(\frac{m \cdot (n + \Delta n) \cdot \Delta n}{|\mathcal{P}|}\right)$ . Once PFVP identifies the violating tuple pairs  $\Delta C_\varphi^p$ , the algorithm IPED scans each tuple pair  $(t_i, t_j)$  in constant time to detect erroneous cells, taking  $O((n + \Delta n) \cdot \Delta n)$  time overall for that final step.

Hence, each constraint  $\varphi \in \Sigma$  requires  $O\left(\frac{m \cdot (n + \Delta n) \cdot \Delta n}{|\mathcal{P}|}\right)$  time for incremental detection, and thus the total complexity over all  $|\Sigma|$  constraints is  $O\left(\frac{|\Sigma| \cdot m \cdot (n + \Delta n) \cdot \Delta n}{|\mathcal{P}|}\right)$  as stated in proposition 2.

